

# CS50: Introduction to Artificial Intelligence Notes

Noah Peters

May 11, 2023

## **Abstract**

CS50: Introduction to Artificial Intelligence notes. Note template by Pingbang Hu.

# Contents

<b>1</b>	<b>Knowledge</b>	<b>2</b>
1.1	Propositional Logic . . . . .	2
1.2	Inference . . . . .	2
1.3	Inference By Resolution . . . . .	3
<b>2</b>	<b>Probability</b>	<b>4</b>
2.1	Introduction to Probability . . . . .	4
2.2	Random Variables . . . . .	4
2.3	Baye's Rule . . . . .	4
2.4	Bayesian Networks . . . . .	5
2.5	Inference by Enumeration . . . . .	6
2.6	Sampling . . . . .	6
2.7	Markov Models . . . . .	6
<b>3</b>	<b>Optimization</b>	<b>8</b>
3.1	Hill Climbing . . . . .	8
3.2	Simulated Annealing . . . . .	8
3.3	Linear Programming . . . . .	9
3.4	Constraint Satisfaction . . . . .	9
3.5	Arc Consistency . . . . .	11
3.6	Backtracking . . . . .	11

# Chapter 1

## Knowledge

### 1.1 Propositional Logic

We use standard logic notation:

- $\neg p$
- $p \vee q$
- $p \wedge q$
- $p \Rightarrow q$ :

$p$	$q$	$p \Rightarrow q$
false	false	true
false	true	true
true	false	false
true	true	true

- $p \Leftrightarrow q$ :

$p$	$q$	$p \Leftrightarrow q$
false	false	true
false	true	false
true	false	false
true	true	true

Now we must establish *what* is considered to be "true" in our world by defining a **model**. We need to represent that knowledge. We do so by defining it via a **knowledge base**.

**Definition 1.1.1 (Model).** Assignment of a truth value to every propositional symbol.

**Definition 1.1.2 (Knowledge Base).** A set of sentences known by a knowledge-based agent.

**Definition 1.1.3 (Entailment).**

$$\alpha \models \beta \text{ "}\alpha \text{ entails } \beta\text{"}$$

In every model in which sentence  $\alpha$  is true, sentence  $\beta$  is also true.

### 1.2 Inference

Our aim is to see if our knowledge base,  $KB$ , entails some query about the world,  $\alpha$ :

$$KB \models \alpha?$$

We first define a **model checking algorithm** to determine if  $KB \models \alpha$ . We can determine this by doing the following:

- enumerate all possible models
- if in every model where  $KB$  is true,  $\alpha$  is also true, then  $KB \models \alpha$

### 1.3 Inference By Resolution

To determine if  $KB \models \alpha$  via knowledge resolution:

- Check if  $KB \wedge \neg\alpha$  is a contradiction
  - Convert  $KB \wedge \neg\alpha$  to Conjunctive Normal Form
  - Keep checking to see if we can use resolution to produce new clause
  - If we ever produce the empty clause (equivalent to False), we have a contradiction and so  $KB \models \alpha$
- If so, then  $KB \models \alpha$
- Otherwise, no entailment

**Problem 1.3.1.** Does  $(A \vee B) \wedge (\neg B \vee C) \wedge (\neg C)$  entail  $A$  ?

**Answer.** First, we convert to CNF:

$$(A \vee B) \wedge (\neg B \vee C) \wedge (\neg C) \wedge (\neg A)$$

We can resolve  $(\neg B \vee C)$  and  $(\neg C)$  by concluding that  $\neg B$ . With the knowledge of  $\neg B$  we now see that, considering  $A \vee B$ , we can conclude  $A$ . We see that

$$A \wedge \neg A \Rightarrow \text{False}$$

and so we can conclude that the clause entails  $A$ . ⊗

# Chapter 2

## Probability

### 2.1 Introduction to Probability

We can represent a **possible world** using  $\omega$  where all possible worlds is the set  $\Omega$ . Thus, we can define total probability as:

$$\sum_{\omega \in \Omega} p(\omega).$$

We build machine models to predict outcomes based on data using **conditional probability**. In other words, we have probability based on some evidence. We can calculate conditional probability using the following formula:

$$P(a|b) = \frac{P(a \wedge b)}{P(b)}.$$

### 2.2 Random Variables

In probability we define a variable with a set of possible values as a **random variable**:

**Definition 2.2.1 (Random Variable).** A variable possessing a distribution of probabilities for various "states".

### 2.3 Baye's Rule

From the above equation we have:

$$P(a \wedge b) = P(b)P(a|b) = P(a)P(b|a)$$

**Definition 2.3.1 (Independence).** The knowledge that one event occurs does not affect the probability of the other event. So we have

$$P(a \wedge b) = P(a)P(b)$$

since  $P(b) = P(b|a)$  if  $a$  and  $b$  are independent.

Thus, we can derive Baye's rule, which relates the probability of one event on the condition of another event, to the reverse relationship.

**Definition 2.3.2 (Baye's Rule).**

$$P(a|b) = \frac{P(b|a)P(a)}{P(b)}$$

What Baye's rule allows us to do is that given:

$$P(\text{visible effect} \mid \text{unknown cause})$$

we can calculate

$$P(\text{unknown cause} \mid \text{visible effect}).$$

## 2.4 Bayesian Networks

There are a number of different probabilistic models. The first we discuss are **Bayesian Networks**.

**Definition 2.4.1 (Bayesian Network).** A data structure that represents the dependencies among random variables. They have the following characteristics:

- directed graph
- each node represents a random variable
- arrow from  $X$  to  $Y$  means  $X$  is a parent of  $Y$
- each node  $X$  has probability distribution  $P(X|\text{Parents}(X))$

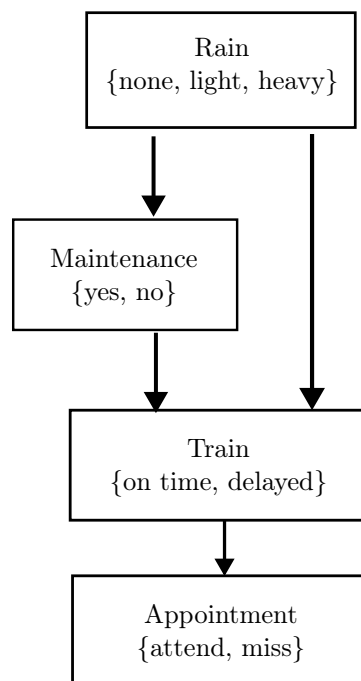


Figure 2.1: A basic example of a Bayesian Network.

We now aim to make an *inference* using the Bayesian Network. Given the following:

- Query  $X$ : variable for which to compute distribution
- Evidence variables  $E$ : observed variables for event  $e$
- Hidden variables  $Y$ : non-evidence, non-query variable

our goal is to calculate  $P(X|e)$ .

**Problem 2.4.1.** Calculate  $P(\text{Appointment}|\text{light, no})$

**Answer.** Here, the *evidence* is that there is light rain and no maintenance. The *query* is the status of Appointment. The *hidden layer* is the status of the train, since you are not given the train's status and you are not querying it; it's just a confounding variable.

---

We note that

$$P(\text{Appointment}|\text{light, no}) = \alpha P(\text{Appointment, light, no})$$

and by the marginalization technique:

$$= \alpha [P(\text{Appointment, light, no, on time}) + P(\text{Appointment, light, no, delayed})].$$

⊛

**Remark.** The **marginalization technique** essentially states that

$$\alpha P(Q, E) = \alpha [P(Q, E, Y) + P(Q, E, \neg Y)]$$

## 2.5 Inference by Enumeration

The above is an example of **inference by enumeration**. More formally, for the following:

- $X$ : the query variable
- $e$ : the evidence
- $y$ : ranges over values of hidden variables
- $\alpha$ : normalizes the result

we have:

$$P(X|e) = \alpha P(X, e) = \alpha \sum_y P(X, e, y)$$

## 2.6 Sampling

Rather than attempting to calculate an exact probability, we can approximate a probability instead via **sampling**. By randomly generating samples for  $n = 1000$  or  $n = 10000$ , we can get fairly useful results.

**Definition 2.6.1 (Rejection Sampling).** **Rejection sampling** is the process of simulating numerous examples from a distribution while considering only the samples that possess the attributes of the desired query.

**Remark.** Rejection sampling is not particularly effective when the evidence you are looking for is fairly unlikely, since you are rejecting a lot of samples. This is inefficient since you are throwing away a large portion of your samples.

**Definition 2.6.2 (Likelihood Weighting).** Rather than sampling everything, in **likelihood weighting**, we start by *fixing* the values for evidence variables. We then sample the non-evidence variables using conditional probabilities in the Bayesian Network. Finally we weight each sample by its *likelihood*.

## 2.7 Markov Models

As opposed to assigning one random variable to a value, but rather an array of random variables to a value over a timescale. This leads to a lot more data, and so we must make some assumptions.

**Definition 2.7.1 (Markov Assumption).** The assumption that the current state depends only on a finite fixed number of previous states.



---

**Definition 2.7.2 (Markov Chain).** A sequence of random variables where the distribution of each variable follows the Markov assumption.

For example, consider the basic Markov Chain

$$X_t \rightarrow X_{t+1}$$

where  $X_t$  represents the weather today, and  $X_{t+1}$  represents the weather tomorrow. We can construct a **transition model** to describe the relationship as follows:

	$X_{t+1}$ : Sunny	$X_{t+1}$ : Rainy
$X_t$ : Sunny	0.8	0.2
$X_t$ : Rainy	0.3	0.7

The first line, for example, translates to "given today was sunny, tomorrow will be sunny with a probability of 0.8 and rainy with probability 0.2". We can refine our definition of the Markov Model further, by introducing *hidden* states.

**Definition 2.7.3 (Hidden Markov Model).** A Markov model for a system with hidden states that generate some observed event.

**Definition 2.7.4 (Filtering).** Given observations from start until now, calculate distribution for current state.

**Definition 2.7.5 (Prediction).** Given observations from start until now, calculate the distribution for a future state.

**Definition 2.7.6 (Smoothing).** Given observations from start until now, calculate the distribution for a past state.

**Definition 2.7.7 (Most Likely Explanation).** Given observations from start until now, calculate most likely sequence of states.

# Chapter 3

## Optimization

### 3.1 Hill Climbing

The most basic form of optimization we explore is that of the local search.

**Definition 3.1.1 (Local Search).** Search algorithms that maintain a single node and searches by moving to a neighboring node.

We can implement a local searching using **hill climbing** in which we do the following:

- Check values of neighbors of current best value
- If a neighbor has a value closer to the desired value, set that value as current best

---

**Algorithm 3.1:** HillClimb

---

```
Data: problem
1 cur = initial state of problem
2 repeat
3   neighbor = best val neighbor of cur
4   if neighbor not better than current then
5     return cur
6   cur = neighbor
7 until better neighbor does not exist
```

---

The problem with this naive approach (**steepest ascent**) is that we may get stuck at a local extremum rather than the global extremum. Further, we may encounter a "shoulder", where several neighboring values are the same, and get stuck.

### 3.2 Simulated Annealing

Although there are a variety of hill climbing algorithms (stochastic, first-choice, beam, etc.). They have their flaws but the theme is that we never go from a good value to a worse value. Thus, we must overcome local extrema. We can tackle this with **simulated annealing**.

**Definition 3.2.1 (Simulated Annealing).** Simulated annealing is akin to a "cooling" physical process.

- Early on, higher "temperature": more likely to accept neighbors that are worse than current state
- Later on, lower "temperature": less likely to accept neighbors that are worse than current state

---

**Algorithm 3.2:** SimulatedAnnealing

---

**Data:** problem, max  
1 cur = initial state of problem  
2 **for**  $t = 1$  to max **do**  
3      $T = \text{Temperature}(t)$   
4     neighbor = random neighbor of cur  
5      $\Delta E$  = how much better neighbor is than cur  
6     **if**  $\Delta E > 0$  **then**  
7         cur = neighbor  
8     with probability  $e^{\Delta E/T}$ , set cur = neighbor

---

### 3.3 Linear Programming

In the context where we are trying to optimize for some function or when we have real values we are often trying to minimize or maximize a cost function given a variety of constraints. This is where **linear programming** becomes useful.

**Definition 3.3.1 (Linear Programming).** Linear programming problems often entail the following:

- Minimize a cost function  $c_1x_1 + c_2x_2 + \dots + c_nx_n$
- With constraints of form  $a_1x_1 + a_2x_2 + \dots + a_nx_n \leq b$  or of form  $a_1x_1 + a_2x_2 + \dots + a_nx_n = b$
- With bounds for each variable  $l_i \leq x_i \leq u_i$

**Problem 3.3.1.** Two machines  $X_1$  and  $X_2$  which cost \$50/hr and \$80/hr to run, respectively. We have the following constraints:

- $X_1$  requires 5 units of labor,  $X_2$  requires 2 units of labor per hour; we have total of 20 units of labor to spend; goal is to minimize cost
- $X_1$  produces 10 units of output per hour,  $X_2$  produces 12 units of output per hour; company needs 90 units of output

**Answer.** We can create a cost function:

$$50x_1 + 80x_2$$

and our constraints:

$$5x_1 + 2x_2 \leq 20$$

$$10x_1 + 12x_2 \geq 90$$

or equivalently:

$$-10x_1 + -12x_2 \leq -90$$

and then solve using a standard linear programming technique such as Simplex, Interior-Point, etc.

⊛

### 3.4 Constraint Satisfaction

Constraint satisfaction problems often have some number of variables that must be optimized, but they are subject to some constraints.


**Definition 3.4.1 (Constraint Satisfaction Problem).** A general constraint satisfaction problem consists of the following:

- Set of variables  $\{x_1, \dots, x_n\}$
- Set of domains for each variable  $\{D_1, \dots, D_n\}$
- Set of constraints  $C$

**Problem 3.4.1.** We must schedule exams for classes  $A, B, \dots, G$  such that no conflicts arise for the students taking the courses. The four students' schedules are:

- *Student 1:*  $A, B, C$
- *Student 2:*  $B, D, E$
- *Student 3:*  $C, E, F$
- *Student 4:*  $E, F, G$

We can represent this graphically using an undirected graph, where the nodes represent exams and the edges indicate that the two exams *cannot* be scheduled for the same time.



Here we define our variables, domains, and constraints:

- Variables:  $\{A, B, C, D, E, F, G\}$
- Domains:  $\{Monday, Tuesday, Wednesday\}$
- Constraints:  $\{A \neq B, A \neq C, B \neq C, \dots, E \neq G, F \neq G\}$

Our goal is to find an assignment of a day to each of the classes such that we don't have any conflicts between the classes. In other words, we are aiming for **node consistency**.

**Definition 3.4.2 (Node Consistency).** When all the values in a variable's domain satisfy the variable's *unary* constraints.

Furthermore, we are seeking **arc consistency**.

**Definition 3.4.3.** When all the values in a variable's domain satisfy the variable's binary constraints. More formally: to make  $X$  arc-consistent with respect to  $Y$ , remove elements from  $X$ 's domain until every choice for  $X$  has a possible choice for  $Y$ .

**Definition 3.4.4 (Unary Constraint).** A constraint involving a *single* variable. (e.g.  $\{A \neq Wednesday\}$ )

**Definition 3.4.5 (Binary Constraint).** A constraint involving *two* variables. (e.g.  $\{A \neq B\}$ )

Note that constraints can come of different forms. **Hard** constraints are absolute requirements, whereas **soft** constraints are preferences.

---

## 3.5 Arc Consistency

We first aim to define a function that, given some consistency problem *csp*, can make some variable *X* **arc consistent** with respect to another variable, *Y*.

---

**Algorithm 3.3:** Revise

---

**Data:** *csp*, *X*, *Y*  
1 revised = false  
2 **for** *x* in *X.domain* **do**  
3     **if** no *y* in *Y.domain* satisfies constraint for (*X*, *Y*) **then**  
4         delete *x* from *X.domain*  
5         revised = true  
6 **return** revised

---

We can enforce arc consistency across an entire consistency problem:

---

**Algorithm 3.4:** AC-3

---

**Data:** *csp*  
1 queue = all arcs in *csp*  
2 **while** queue not empty **do**  
3     (*X*, *Y*) = Dequeue(queue)  
4     **if** Revise(*csp*, *X*, *Y*) **then**  
5         **if** size of *X.domain* == 0 **then**  
6             **return** false  
7         **for** each *Z* in *X.neighbors* - {*Y*} **do**  
8             Enqueue(queue, (*Z*, *X*))  
9 **return** true

---

### 3.5.1 CSPs as Search Problems

We can reframe CSPs as *search problems*. We do so by defining a CSP as follows:

- initial state: empty assignment (no variables)
- actions: add a {*variable* = *value*} to assignment
- transition model: shows how adding an assignment changes the assignment
- goal test: check if all variables assigned and constraints all satisfied
- path cost function: all paths have same cost

## 3.6 Backtracking

We can find a solution to a CSP by simply applying arbitrary assignments to variables one by one until a constraint is broken. Then, we can just backtrack and try another assignment. We eventually find a solution or check every possible assignment.

---

**Algorithm 3.5:** Backtrack

---

**Data:** assignment, csp

```
1 if assignment complete then
2   return assignment
3 var = SelectUnassignedVar(assignment, csp)
4 for value in DomainValues(var, assignment, csp) do
5   if value consistent with assignment then
6     add var = value to assignment
7     result = Backtrack(assignment, csp)
8     if result  $\neq$  failure then
9       return result
10  remove var = value from assignment
11 return failure
```

---

# Appendix