

# EECS 370: Introduction to Computer Organization Notes

Noah Peters

March 28, 2023

## **Abstract**

Lecture notes for EECS 370 at the University of Michigan. L<sup>A</sup>T<sub>E</sub>Xtemplate by Pingbang Hu.

# Contents

<b>3</b>	<b>Caches</b>	<b>2</b>
3.1	Memory Hierarchy . . . . .	2
3.2	Cache Basics . . . . .	3
3.3	Cache Organization . . . . .	5
<b>A</b>	<b>Additional Proofs</b>	<b>7</b>
A.1	Proof of ?? . . . . .	7

# Chapter 3

## Caches

### Lecture 17: Introduction to Caches

#### 3.1 Memory Hierarchy

14 Mar. 12:00

We often need a lot of memory, LC2K alone can handle  $2^{18}$  bytes of memory. We have several choices for memory:

- **SRAM:** Static Random Access Memory
  - fast: 2ns access time or faster
  - decoders are big
  - expensive, high area requirement
- **DRAM:** Dynamic Random Access Memory
  - slower: 50ns access time
  - must stall for dozens of cycles each memory load
  - less expensive
- **Flash**
  - slow: access time varies wildly
  - less expensive than DRAM
  - non-volatile
- **Disks**
  - obnoxiously slow: 3,000,000ns access time
  - dirt cheap
  - non-volatile

As seen above, there are trade-offs among each type of memory. Ideally, we would have cheap *and* fast memory. So, we can use a combination of memory types to optimize the common case via strategic **locality of reference**.




Figure 3.1: Caches Overview

**Definition 3.1.1.** The **architectural** view of memory is what the machine language (or programmer) sees, i.e. just a big array.

**Note.** Breaking up the memory system into different pieces (cache, main memory, disk) is **not architectural**. The machine language doesn't know about it.

Can use our variety of memories as follows:

- use small array of SRAM for the **cache**
- use a larger amount of DRAM for **main memory**
- use a lot of flash and/or disk for **virtual memory**

## 3.2 Cache Basics

Whenever memory returns data, we can store it in a cache. We'll need to store:

- the data
- a tag denoting its memory location
- a "valid" status bit

Then for our next memory access, we can first check if the tag matches the address we are attempting to access.

**Definition 3.2.1 (Cache Hit).** A **hit** occurs when data for a memory access is found in the cache.

**Definition 3.2.2 (Cache Miss).** A **miss** occurs when data for a memory access is *not* found in the cache.

**Definition 3.2.3 (Hit/Miss Rate).** The **hit/miss rate** is the percentage of memory accesses that hit/miss in the cache.

### 3.2.1 CAMs

**Definition 3.2.4. CAMs: content addressable memories** are akin to a set of data matching a query. Instead of an address we send a key to the memory, asking whether the key exists and, if so, what value it is associated with. Memory answers: yes/no and gives associated value (if there is one).

We apply *operations* on CAMS:

- **Search:** the primary way to access a CAM
  - send data to CAM memory
  - return "found" or "not found"
  - if found, return location of where it was found or its associated value
- **Write:** send data for CAM to remember


### 3.2.2 Cache Organization

Cache memory can copy data from any part of main memory. Cache memory has two parts:

- the **tag (CAM)**: holds the memory address
- the **block (SRAM)**: holds the memory data

A **hit** in the cache occurs when a tag match is found. The microprocessor sends the address to the CAM containing the tags and searches for the tag. If there's a search result hit, the corresponding block is forwarded to the microprocessor. If not, the address is forwarded to main memory:

Hit:



Miss:

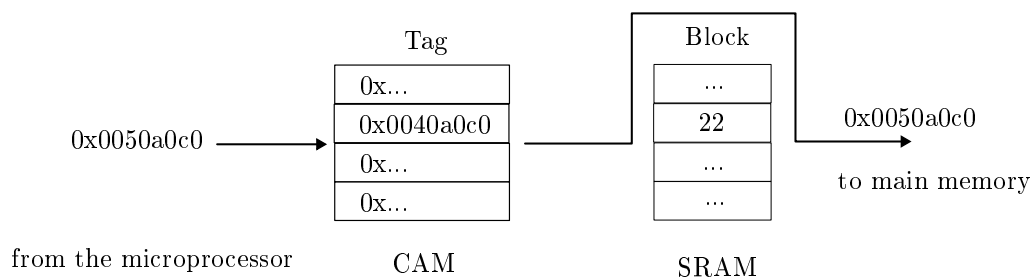


Figure 3.2: Hardware view of caches.

**Problem 3.2.1.** Given:

- *cache* has 1 cycle access time
- *main memory* has 100 cycle access time
- *disk* has 10,000 cycles access time

What is the average access time for 100 memory references if 90% of the cache accesses are hits

---

and 80% of the accesses to main memory are hits? Assume main memory access time includes tag array access to determine hit/miss.

**Answer.**  $0.9 \cdot 1 + 0.1 \cdot (100 + 0.2 \cdot 10000) = 210.9$  ⊗

### 3.2.3 Cache Operation

Every cache *miss* will get the data from memory and *allocate* a cache line to put the data in (just like any CAM write). However... what do we replace in the cache? Does an optimal replacement policy exist?

**Definition 3.2.5 (Temporal locality).** The principle of **temporal locality** in program references says that if you access a memory location (e.g., 1000) you will be more likely to re-access that location than you will be to reference some other random location.

**Remark.** Locality is a property of *programs* (not hardware).

Specifically, temporal locality says that the **least recently referenced (LRU)** cache line should be *evicted* to make room for the new line. Because the re-access probability falls over time as a cache line isn't referenced, the LRU line is least likely to be re-referenced.

**Definition 3.2.6 (Average Access Latency).** Average Access Latency = cache latency · hit rate + memory latency ·

## Lecture 18: Cache Organization: Block Size and Writes

### 3.3 Cache Organization

16 Mar. 12:00

**Definition 3.3.1 (k-Way Set Associative LRU).** A set associative LRU with  $k$  **ways**, has essentially  $k$  cells within each cache line's block. The LRU is set to be the *way* with the highest *count*. 20:43

# Appendix



# Appendix A

## Additional Proofs

### A.1 Proof of ??

We can now prove ??.

**Proof of ??.** See [here](#).

